

# Machine Learning Analysis of Titanic

## 1. Import Necessary Library

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import sklearn
import seaborn as sns
```

## 2. Data Preprocessing

### 2.1 Preview: Overround Understanding of Data

```
In [ ]: # use two copy of datas
# the library's file outlines 'read_excel' output is DataFrame type
data = pd.read_excel('titanic.xlsx')
raw_data = pd.read_excel('titanic.xlsx')

data.head()
```

```
Out[ ]: 
```

	pclass	survived	name	sex	age	sibsp	parch	ticket	fare	ca
0	1	1	Allen, Miss. Elisabeth Walton	female	29.0000	0	0	24160	211.3375	
1	1	1	Allison, Master. Hudson Trevor	male	0.9167	1	2	113781	151.5500	(
2	1	0	Allison, Miss. Helen Loraine	female	2.0000	1	2	113781	151.5500	(
3	1	0	Allison, Mr. Hudson Joshua Creighton	male	30.0000	1	2	113781	151.5500	(
4	1	0	Allison, Mrs. Hudson J C (Bessie Waldo Daniels)	female	25.0000	1	2	113781	151.5500	(

view data's stat info & description

## 2.2 More Detailed: Clear Discription of Data (Types/Null)

```
In [ ]: print(f'Data Shape: {data.shape}\n')
        print(f'Columns: {data.columns}\n')
        print(f'Info: {data.info()}\n')
        print(f'Describe: {data.describe()}')
```

Data Shape: (1309, 14)

Columns: Index(['pclass', 'survived', 'name', 'sex', 'age', 'sibsp', 'parch', 'ticket', 'fare', 'cabin', 'embarked', 'boat', 'body', 'home.dest'], dtype='object')

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1309 entries, 0 to 1308

Data columns (total 14 columns):

#	Column	Non-Null Count	Dtype
0	pclass	1309 non-null	int64
1	survived	1309 non-null	int64
2	name	1309 non-null	object
3	sex	1309 non-null	object
4	age	1046 non-null	float64
5	sibsp	1309 non-null	int64
6	parch	1309 non-null	int64
7	ticket	1309 non-null	object
8	fare	1308 non-null	float64
9	cabin	295 non-null	object
10	embarked	1307 non-null	object
11	boat	486 non-null	object
12	body	121 non-null	float64
13	home.dest	745 non-null	object

dtypes: float64(3), int64(4), object(7)

memory usage: 143.3+ KB

Info: None

Describe:	pclass	survived	age	sibsp
parch \				
count	1309.000000	1309.000000	1046.000000	1309.000000
mean	2.294882	0.381971	29.881135	0.385027
std	0.837836	0.486055	14.413500	0.865560
min	1.000000	0.000000	0.166700	0.000000
25%	2.000000	0.000000	21.000000	0.000000
50%	3.000000	0.000000	28.000000	0.000000
75%	3.000000	1.000000	39.000000	0.000000
max	3.000000	1.000000	80.000000	9.000000

	fare	body
count	1308.000000	121.000000
mean	33.295479	160.809917
std	51.758668	97.696922
min	0.000000	1.000000
25%	7.895800	72.000000
50%	14.454200	155.000000
75%	31.275000	256.000000
max	512.329200	328.000000

we find that the column `Non-Null Count` assigned: **if the Non-Null Count != `data.shape[0]`**, it means this feature attribute column **has null value**. like the following columns have null-values, some even has a large amount of null value:

- 4	age	1046 non-null	float64
- 8	fare	1308 non-null	float64
- 9	cabin	295 non-null	object
- 10	embarked	1307 non-null	object
- 11	boat	486 non-null	object
- 12	body	121 non-null	float64
- 13	home.dest	745 non-null	object

However, although we know some of the columns have null-values, **but we dont directly cope with these null values**, instead, we make a simple visualization for these data.

## 2.3 Visualization: Count Visualization of Numerical / Categorical Data

```
In [ ]: sns.set(style="whitegrid")

fig, axes = plt.subplots(4, 5, figsize=(20, 15))

# age
sns.histplot(data['age'].dropna(), kde=True, bins=30, color='skyblue', ax=axes[0, 0])
axes[0, 0].set_title('Age Distribution')
axes[0, 0].set_xlabel('Age')
axes[0, 0].set_ylabel('Frequency')

# sibps
sns.histplot(data['sibsp'], kde=False, bins=10, color='salmon', ax=axes[0, 1])
axes[0, 1].set_title('Siblings/Spouses Aboard Distribution')
axes[0, 1].set_xlabel('Number of Siblings/Spouses')
axes[0, 1].set_ylabel('Frequency')

# parch
sns.histplot(data['parch'], kde=False, bins=10, color='green', ax=axes[0, 2])
axes[0, 2].set_title('Parents/Children Aboard Distribution')
axes[0, 2].set_xlabel('Number of Parents/Children')
axes[0, 2].set_ylabel('Frequency')

# fare
sns.histplot(data['fare'], kde=True, bins=30, color='purple', ax=axes[0, 3])
axes[0, 3].set_title('Fare Distribution')
axes[0, 3].set_xlabel('Fare')
axes[0, 3].set_ylabel('Frequency')

# body
sns.histplot(data['body'].dropna(), kde=True, bins=20, color='orange', ax=axes[1, 0])
axes[1, 0].set_title('Body Condition Distribution')
axes[1, 0].set_xlabel('Body Condition')
axes[1, 0].set_ylabel('Frequency')

# pclass
sns.countplot(x='pclass', data=data, palette='pastel', ax=axes[1, 1])
```

```

axes[1, 1].set_title('Passenger Class Distribution')
axes[1, 1].set_xlabel('Pclass')
axes[1, 1].set_ylabel('Count')

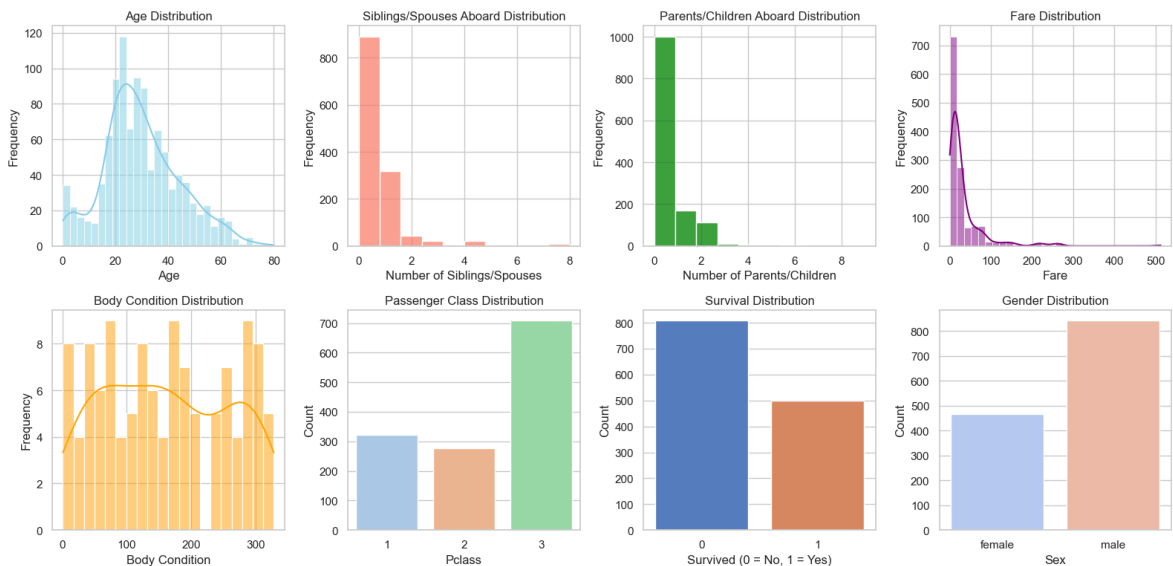
# survived
sns.countplot(x='survived', data=data, palette='muted', ax=axes[1, 2])
axes[1, 2].set_title('Survival Distribution')
axes[1, 2].set_xlabel('Survived (0 = No, 1 = Yes)')
axes[1, 2].set_ylabel('Count')

# sex
sns.countplot(x='sex', data=data, palette='coolwarm', ax=axes[1, 3])
axes[1, 3].set_title('Gender Distribution')
axes[1, 3].set_xlabel('Sex')
axes[1, 3].set_ylabel('Count')

for ax in axes.flat:
    if not ax.has_data():
        ax.set_visible(False)

plt.tight_layout()
plt.show()

```



## 2.4 Analysis on Visualization: Make Explanation of Graph

- **Age** : most of passengers are centred on age: **16~35**
- **Number of Siblings** : most of them is **the only traveler, not accompany with their own brothers or sisters** aboard. (Ops, it maybe a bit sad, cuz the loss take away the joyful and happiness of the whole family, what a pity tbh)
- **Number of Parents** : most of them is **the only traveler, not accompany with their own parents or children**, it means **most of them are couple!**
- **Fare** : most of them pay a low-price ticket to get aboard. (just like the Movie shown, similar as Jack hah)

- **Body Condition** : body weight (/lb), the distribution is balanced and most of them are centered on range 50lb~200+lb
- **Pclass** : the First Class / Second Class / Third Class (amazingly to find that there are more First Class than the Second Class instead)
- **Survived** : the **target value column**, it breaks the previous cognition of my childhood that 'most of the people in Titanic died', cuz the first impression left by the Movie 'The Titanic'. **But according to stat info, near up to 40% of people survived**
- **Sex** : the male is near 65%

## 2.5 Feature Selection and Dimension Reduction

we know the raw data shape is `(1309, 14)`, that's not a huge dataset in fact. But we notice that **some of the features may be not that helpful for model training**. Or **some data's null ratio is high**.

### Feature Selection

We need to select necessary features as follows:

```
# (col drop) drop some columns
data = data.drop(columns = [__columnsToDrop__])

# (row drop) combined condition to judge the row
data = data.dropna(subset=[__columnsToDrop__], how='all')
```

- Drop some **unnecessary or empty-main** columns
- Drop rows: cuz **cabin & age** are both important factor for survival rate. **drop if a single miss both of these important features**

```
In [ ]: print(f'Raw Data Shape: {data.shape}\n')
        print(f'Columns: {data.columns}\n')
        data = data.drop(columns = ['survived', 'name', 'ticket', 'body', 'home.d

        print(f'Dropped Data Shape: {data.shape}\n')
        print(f'Cleaned Data: {data.columns}')

        data = data.dropna(subset=['cabin', 'age'], how='all')
        raw_data = raw_data.dropna(subset=['cabin', 'age'], how='all')

        print('----' * 28)
```

Raw Data Shape: (1309, 14)

```
Columns: Index(['pclass', 'survived', 'name', 'sex', 'age', 'sibsp', 'parch', 'ticket',
               'fare', 'cabin', 'embarked', 'boat', 'body', 'home.dest'],
              dtype='object')
```

Dropped Data Shape: (1309, 8)

```
Cleaned Data: Index(['pclass', 'sex', 'age', 'sibsp', 'parch', 'fare', 'cabin', 'boat'],
                    dtype='object')
```

```
In [ ]: print(data.head(), '\n')
        print(data.shape)
```

	pclass	sex	age	sibsp	parch	fare	cabin	boat
0	1	female	29.0000	0	0	211.3375	B5	2
1	1	male	0.9167	1	2	151.5500	C22 C26	11
2	1	female	2.0000	1	2	151.5500	C22 C26	NaN
3	1	male	30.0000	1	2	151.5500	C22 C26	NaN
4	1	female	25.0000	1	2	151.5500	C22 C26	NaN

(1069, 8)

## 2.6 NULL Value Processing

after data feature selection, we need to find **null value** and process with it.

```
In [ ]: null_col = list()
        container = list()

        for i in list(data.columns):
            # print(i)
            if data[i].isna().any() == True:
                null_col.append(i)

        print(f'Orinal Has Null Col: {null_col}')

        # using mode value to fill up
        for col in list(data.columns):
            # print(data[col].mode())
            # avg = data[col].mode()[0]
            mode = data[col].mode()[0]
            data[col] = data[col].fillna(mode)

        for i in list(data.columns):
            # print(i)
            if data[i].isna().any() == True:
                container.append(i)

        print(f'Processed Has Null Col: {container}')
```

Orinal Has Null Col: ['age', 'fare', 'cabin', 'boat']

Processed Has Null Col: []

but sadly, when I check in the data, I found **some error in boat column**, like some input are '14 15 B', 'B'.... It means some **Non-Numeric Data involved Numeric**

Column (Code is following the encoding part)

now, after we fill the NULL value with its `mode-filled` of responding column, we get as follows:

*# means no Null Col, all the Col have no null value*

Processed Has Null Col: []

## 2.7 Encoding of Categorical Data











most of the Categorical Data can't be directly used in Model Learning Model, so  
transform Categorical Data --> Numeric Data

- (1) One Hot-Encoding : For a certain feature, it has many categorical types. For this certain feature, we make Binary Encoding for these types. **Create multiple new columns, each column stands for a certain type of certain feature**
  - the following image, we can see **three new columns were born**, (if we have more categorical features, that means more types. So more new columns will be generated)

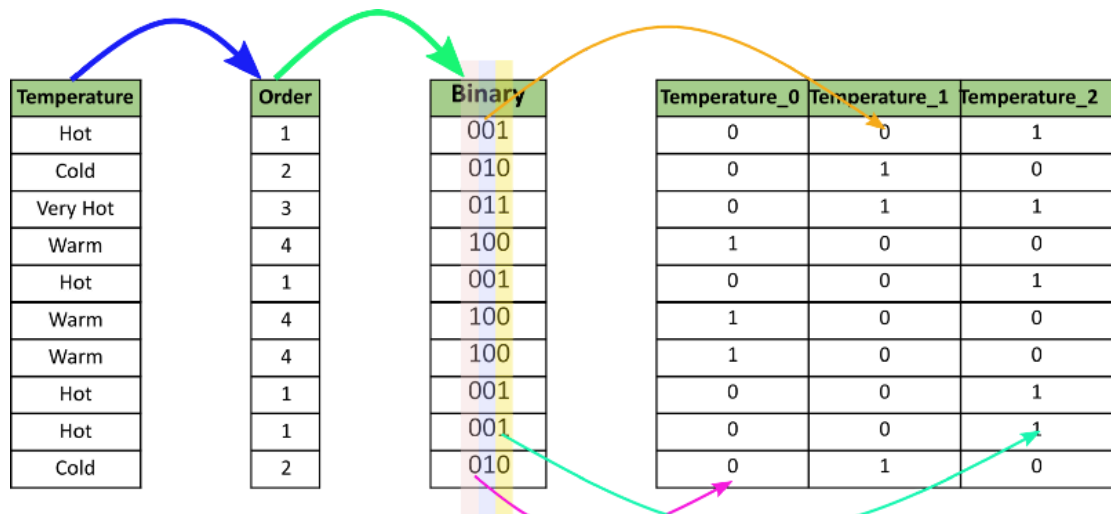
Color		Red	Yellow	Green
Red		1	0	0
Red		1	0	0
Yellow		0	1	0
Green		0	0	1
Yellow				

– understand how to assign a certain categorical type

### One hot encoding

Ear-shape	Pointy ears	Floppy ears	Oval ears	Face shape	Whiskers	Cat
 Pointy	1	0	0	Round	Present	1
 Oval	0	0	1	Not round	Present	1
 Oval	0	0	1	Round	Absent	0
 Pointy	1	0	0	Not round	Present	0
 Oval	0	0	1	Round	Present	1
 Pointy	1	0	0	Round	Absent	1
 Floppy	0	1	0	Not round	Absent	0
 Oval	0	0	1	Round	Absent	1
 Floppy	0	1	0	Round	Absent	0
 Floppy	0	1	0	Round	Absent	0

- (2) Label Encoding : Direct label by **0/1/2/3/4/...**
- (3) Binary Encoding : advanced One Hot-Encoding, can reduce the data dimension



```
In [ ]: # data encoding
# it use one hot-encoding
data_encoded = pd.get_dummies(
    data, columns=['sex', 'cabin']
)

print(data_encoded)
```



	pclass	age	sibsp	parch	fare	boat	sex_female	sex_male
\								
0	1	29.0000	0	0	211.3375	2	True	False
1	1	0.9167	1	2	151.5500	11	False	True
2	1	2.0000	1	2	151.5500	15	True	False
3	1	30.0000	1	2	151.5500	15	False	True
4	1	25.0000	1	2	151.5500	15	True	False
...	...	...	...	...	...	...	...	...
1301	3	45.5000	0	0	7.2250	15	False	True
1304	3	14.5000	1	0	14.4542	15	True	False
1306	3	26.5000	0	0	7.2250	15	False	True
1307	3	27.0000	0	0	7.2250	15	False	True
1308	3	29.0000	0	0	7.8750	15	False	True

	cabin_A10	cabin_A11	...	cabin_F E57	cabin_F E69	cabin_F G63	\
0	False	False	...	False	False	False	
1	False	False	...	False	False	False	
2	False	False	...	False	False	False	
3	False	False	...	False	False	False	
4	False	False	...	False	False	False	
...	...	...	...	...	...	...	
1301	False	False	...	False	False	False	
1304	False	False	...	False	False	False	
1306	False	False	...	False	False	False	
1307	False	False	...	False	False	False	
1308	False	False	...	False	False	False	

	cabin_F G73	cabin_F2	cabin_F33	cabin_F38	cabin_F4	cabin_G6	cab
in_T							
0	False	False	False	False	False	False	F
alse							
1	False	False	False	False	False	False	F
alse							
2	False	False	False	False	False	False	F
alse							
3	False	False	False	False	False	False	F
alse							
4	False	False	False	False	False	False	F
alse							
...	...	...	...	...	...	...	
...							
1301	False	False	False	False	False	False	F
alse							
1304	False	False	False	False	False	False	F
alse							
1306	False	False	False	False	False	False	F
alse							
1307	False	False	False	False	False	False	F
alse							
1308	False	False	False	False	False	False	F
alse							

[1069 rows x 194 columns]

using **One Hot-Encoding** , and found that data dimension is improving! Turn into  
higher dimension : [[1069 rows x 194 columns]]

```
In [ ]: # found new generated columns
data_encoded[['cabin_F2', 'cabin_F33', 'cabin_F38', 'cabin_F4']].head()
```

```
Out [ ]:      cabin_F2  cabin_F33  cabin_F38  cabin_F4
```

	cabin_F2	cabin_F33	cabin_F38	cabin_F4
0	False	False	False	False
1	False	False	False	False
2	False	False	False	False
3	False	False	False	False
4	False	False	False	False

```
In [ ]: # some error in `boat` column, like some input are '14 15 B', 'B'.... It

data_encoded['boat'] = pd.to_numeric(data['boat'], errors='coerce')
# data_encoded['boat'].head(30)

mean1 = data_encoded['boat'].mean()
data_encoded['boat'] = data_encoded['boat'].fillna(mean1)

data_encoded['boat'].head(10)
```

```
Out [ ]: 0      2.000000
1      11.000000
2      15.000000
3      15.000000
4      15.000000
5       3.000000
6      10.000000
7      15.000000
8      12.872128
9      15.000000
Name: boat, dtype: float64
```

## 2.8 Overview of Processed Data

```
In [ ]: # choose the first 6 columns
# cuz the following columns are both new generated columns for One Hot-En
for i in list(data_encoded.columns)[0:6]:
    print(f'{i}: {(set(data[i]))}\n')

# data_encoded.head()
```

pclass: {1, 2, 3}

age: {0.75, 1.0, 2.0, 3.0, 4.0, 5.0, 6.0, 7.0, 8.0, 9.0, 10.0, 11.0, 12.0, 13.0, 14.0, 15.0, 16.0, 17.0, 18.0, 19.0, 18.5, 21.0, 22.0, 23.0, 24.0, 25.0, 26.0, 27.0, 28.5, 28.0, 29.0, 30.0, 31.0, 32.0, 33.0, 32.5, 35.0, 36.0, 37.0, 38.0, 39.0, 40.0, 41.0, 42.0, 43.0, 44.0, 45.0, 46.0, 47.0, 48.0, 49.0, 50.0, 51.0, 52.0, 53.0, 54.0, 55.0, 56.0, 57.0, 58.0, 59.0, 60.0, 61.0, 62.0, 63.0, 64.0, 65.0, 66.0, 67.0, 60.5, 70.0, 71.0, 70.5, 14.5, 74.0, 76.0, 80.0, 20.0, 20.5, 22.5, 23.5, 24.5, 0.3333, 26.5, 30.5, 34.0, 34.5, 36.5, 38.5, 40.5, 45.5, 0.9167, 0.8333, 0.6667, 55.5, 11.5, 0.1667, 0.4167}

sibsp: {0, 1, 2, 3, 4, 5, 8}

parch: {0, 1, 2, 3, 4, 5, 6}

fare: {0.0, 512.3292, 3.1708, 4.0125, 5.0, 6.75, 7.55, 7.65, 9.6875, 10.5, 11.5, 12.525, 13.0, 13.5, 13.8583, 14.5, 16.0, 12.275, 15.0, 13.7917, 21.0, 15.0458, 23.0, 24.0, 25.5875, 26.3875, 27.75, 28.5, 28.7125, 26.0, 26.5, 27.7208, 26.2875, 30.0, 31.0, 30.5, 29.7, 35.5, 31.6792, 38.5, 39.6, 34.6542, 42.5, 39.4, 45.5, 42.4, 47.1, 40.125, 49.5042, 50.4958, 51.8625, 51.4792, 52.5542, 53.1, 52.0, 55.0, 56.9292, 57.0, 59.4, 55.4417, 61.175, 57.9792, 63.3583, 57.75, 61.9792, 66.6, 60.0, 61.3792, 69.3, 65.0, 71.0, 71.2833, 73.5, 14.0, 75.2417, 76.2917, 77.9583, 78.85, 78.2667, 79.2, 81.85, 83, 76.7292, 83.1583, 83.475, 80.0, 86.5, 82.1708, 82.2667, 89.1042, 90.0, 91.0792, 18.75, 93.5, 18.7875, 17.4, 19.5, 19.2583, 18.0, 7.2292, 20.25, 7.8542, 16.1, 20.525, 20.2125, 106.425, 108.9, 110.8833, 22.525, 22.025, 13.275, 22.3583, 120.0, 24.15, 25.9292, 26.25, 14.1083, 133.65, 134.5, 135.6333, 136.7792, 27.4458, 27.0, 8.6833, 27.9, 28.5375, 146.5208, 29.0, 29.125, 151.55, 30.6958, 153.4625, 30.0708, 31.5, 31.275, 6.4375, 32.3208, 32.5, 164.8667, 7.775, 7.8208, 33.5, 33.0, 8.9625, 10.1708, 34.0208, 34.375, 9.5875, 9.8375, 35.0, 10.4625, 7.25, 36.75, 7.75, 11.1333, 7.125, 7.0, 37.0042, 7.875, 12.65, 8.4042, 13.9, 13.4167, 13.775, 39.0, 39.6875, 14.4, 15.9, 41.5792, 211.3375, 211.5, 221.7792, 6.2375, 227.525, 7.925, 7.05, 7.8792, 7.2833, 7.6292, 9.5, 46.9, 7.8, 9.0, 49.5, 247.5208, 50.0, 262.375, 263.0, 55.9, 56.4958, 8.05, 8.3, 8.1583, 8.3625, 8.6542, 9.8458, 9.825, 8.4333, 8.0292, 12.0, 12.875, 9.4833, 9.2167, 9.325, 9.8417, 12.7375, 12.183, 13.8625, 15.0333, 14.4542, 15.2458, 15.55, 16.7, 17.8, 69.55, 6.95, 6.45, 6.4958, 7.8875, 7.4958, 20.575, 21.075, 75.25, 15.75, 25.7, 15.5, 26.2833, 77.2875, 79.65, 7.8292, 7.5792, 7.1417, 31.3875, 6.975, 7.225, 7.725, 7.7417, 7.85, 7.0542, 8.6625, 8.85, 9.475, 9.35, 9.225, 8.5167, 11.2417, 12.35, 12.475, 12.2875, 14.4583, 15.85, 15.1, 15.7417, 7.7958, 7.0458, 7.7333, 10.5167, 7.8958, 7.5208}

boat: {1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 'C', 13, 14, 15, 12, 16, '13 15', 'D', 'C D', 'B', 'A', '8 10', '5 7', '13 15 B', '5 9'}

luckily, we all successfully cope with data preprocessing, now no Null value or illegal value exist!

## 2.9 Data Scaling

for a better control in Gradient Descent, (for a better grad value computation and a more balanced grad leading guidance). **Data Scaling has better control & convergence in gradient computation (faster + stable + more balanced)**

Tips:

- 1. Scaling used in **Numeric** Data
- 2. Scaling used in **X**, but **not y**

```
In [ ]: from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
data_standardized = data_encoded.copy()

# choosing numeric columns,
# the target column doesn't need to!
data_standardized[['age', 'sibsp', 'parch', 'fare', 'boat']] = scaler.fit

print(data_standardized.shape)
data_standardized.head()
```

(1069, 194)

```
Out [ ]:
```

	pclass	age	sibsp	parch	fare	boat	sex_female	sex
0	1	-0.052857	-0.548962	-0.496538	3.140614	-2.943948	True	
1	1	-2.019983	0.556199	1.905267	2.064310	-0.506934	False	
2	1	-1.944102	0.556199	1.905267	2.064310	0.576184	True	
3	1	0.017189	0.556199	1.905267	2.064310	0.576184	False	
4	1	-0.333041	0.556199	1.905267	2.064310	0.576184	True	

5 rows × 194 columns

## 3. Model Preparation

### 3.1 Dataset Split

```
In [ ]: # Split data
from sklearn.model_selection import train_test_split

print(f'Raw Data Shape: {data.shape}')
X_train, X_test, y_train, y_test = train_test_split(
    data_standardized,
    raw_data['survived'],
    test_size = 0.2, random_state = 42
)

print(f'X_train.shape: {X_train.shape}')
print(f'y_train.shape: {y_train.shape}')
print(f'X_test.shape: {X_test.shape}')
print(f'y_test.shape: {y_test.shape}')
```

Raw Data Shape: (1069, 8)  
X\_train.shape: (855, 194)  
y\_train.shape: (855,)  
X\_test.shape: (214, 194)  
y\_test.shape: (214,)

## 3.2 Model Initialization

Simple Binary Classification Problem --> Logistic Model

- Logistic Model with L1 Regularization
- Logistic Model with L2 Regularization

```
In [ ]: from sklearn.linear_model import LogisticRegression
        from sklearn.metrics import accuracy_score

        # Using L1 Regularization
        # init the model
        model_l1 = LogisticRegression(penalty='l1', solver='saga', max_iter=5000,

#####

        # Using L2 regularization
        # init the model
        model_l2 = LogisticRegression(penalty='l2', solver='lbfgs', max_iter=5000

        print(f'Model1: {model_l1}\n')
        print(f'Model2: {model_l2}')
```

Model1: LogisticRegression(max\_iter=5000, penalty='l1', random\_state=42, solver='saga')

Model2: LogisticRegression(max\_iter=5000, random\_state=42)

## 3.3 Model Training

```
In [ ]: # training the Model
        print(model_l1.fit(X_train, y_train))
        print(model_l2.fit(X_train, y_train))
        model_l1
```

LogisticRegression(max\_iter=5000, penalty='l1', random\_state=42, solver='saga')

LogisticRegression(max\_iter=5000, random\_state=42)

```
Out [ ]: ▼ LogisticRegression
          LogisticRegression(max_iter=5000, penalty='l1', random_state=42,
                             solver='saga')
```

## 3.4 Model Prediction

```
In [ ]: y_pred_l1 = model_l1.predict(X_test)
        y_pred_l2 = model_l2.predict(X_test)

        # show one of pred result
        y_pred_l1
```

```
Out[ ]: array([0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0,
        0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0,
        0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0,
        1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0,
        1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1,
        0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
        0, 1, 1, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 1,
        1, 1, 1, 0, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0,
        0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1])
```

### 3.5 Model Evaluation (Quantitative Statistics)

```
In [ ]: # acr print out
# Calculate the accuracy of the model
accuracy_l1 = accuracy_score(y_test, y_pred_l1)
print(f"-- Accuracy with L1 regularization: {accuracy_l1*100:.4f}%")

accuracy_l2 = accuracy_score(y_test, y_pred_l2)
print(f"-- Accuracy with L2 regularization: {accuracy_l2*100:.4f}%\n")

print('---'*18)
# Calculate other metrics: precision, recall, and F1-score
from sklearn.metrics import precision_score, recall_score, f1_score

prec_l1 = precision_score(y_test, y_pred_l1)
recall_l1 = recall_score(y_test, y_pred_l1)
f1_l1 = f1_score(y_test, y_pred_l1)

prec_l2 = precision_score(y_test, y_pred_l2)
recall_l2 = recall_score(y_test, y_pred_l2)
f1_l2 = f1_score(y_test, y_pred_l2)

print("\nL1 Regularization Model:")
print(f"-- Precision: {prec_l1:.4f}")
print(f"-- Recall: {recall_l1:.4f}")
print(f"-- F1 Score: {f1_l1:.4f}\n")

print("L2 Regularization Model:")
print(f"-- Precision: {prec_l2:.4f}")
print(f"-- Recall: {recall_l2:.4f}")
print(f"-- F1 Score: {f1_l2:.4f}\n")

print('---'*18)

# print out classification report
from sklearn.metrics import classification_report

print('\nL1 Regularization Model Report:\n', classification_report(
    y_test, y_pred_l1,
    target_names=['Survived', 'Death'])
)

print('---' * 18)
print('\nL2 Regularization Model Report:\n', classification_report(
    y_test, y_pred_l2,
    target_names=['Survived', 'Death'])
)
```

- Accuracy with L1 regularization: 92.5234%
- Accuracy with L2 regularization: 92.0561%

---

#### L1 Regularization Model:

- Precision: 0.9870
- Recall: 0.8352
- F1 Score: 0.9048

#### L2 Regularization Model:

- Precision: 0.9868
  - Recall: 0.8242
  - F1 Score: 0.8982
- 

#### L1 Regularization Model Report:

	precision	recall	f1-score	support
Survived	0.89	0.99	0.94	123
Death	0.99	0.84	0.90	91
accuracy			0.93	214
macro avg	0.94	0.91	0.92	214
weighted avg	0.93	0.93	0.92	214

---

#### L2 Regularization Model Report:

	precision	recall	f1-score	support
Survived	0.88	0.99	0.93	123
Death	0.99	0.82	0.90	91
accuracy			0.92	214
macro avg	0.94	0.91	0.92	214
weighted avg	0.93	0.92	0.92	214

## 3.6 Model Evaluation (Visualization)

### Plot 1: Confusion Matrix

```
In [ ]: # Generate a confusion matrix and visualize it
from sklearn.metrics import confusion_matrix

# compute the matrix
matrixs_l1 = confusion_matrix(y_test, y_pred_l1)
matrixs_l2 = confusion_matrix(y_test, y_pred_l2)

fig, axes = plt.subplots(1, 2, figsize=(15, 6))


# sub plotting
sns.heatmap(matrixs_l1, annot=True, fmt='g', cmap='Blues', ax=axes[0], xt
axes[0].set_title('L1 Regularization Confusion Matrix')
axes[0].set_xlabel('Predicted Labels')
axes[0].set_ylabel('True Labels')
```

```

sns.heatmap(matrixs_l2, annot=True, fmt='g', cmap='Reds', ax=axes[1], xti
axes[1].set_title('L2 Regularization Confusion Matrix')
axes[1].set_xlabel('Predicted Labels')
axes[1].set_ylabel('True Labels')

plt.tight_layout()
plt.show()


```



## Plot 2: ROC Curve

### About ROC

- the curve **more near upper left**, the **better model performs**
- AUC** : Area Under the Curve, the value determines good / bad the model performs (positive relationship)
  - AUC = 0.5: model preforms random prediction simulation
  - AUC = 1: the most perfect model (**threshold level**)
  - AUC < 0.5: worst than random prediction
- Advantage: fit for unbalanced data distribution (performs well in class imbalance)





```

In [ ]: # Plot the ROC curve and calculate the AUC
        from sklearn.metrics import roc_curve, auc

        # prob transformation
        y_pred_l1_prob = model_l1.predict_proba(X_test)[: , 1]
        print(model_l1.predict_proba(X_test)[:3])

        y_pred_l2_prob = model_l2.predict_proba(X_test)[: , 1]
        print('\n', model_l2.predict_proba(X_test)[:3])

        fpr_l1, tpr_l1, _ = roc_curve(y_test, y_pred_l1_prob)
        fpr_l2, tpr_l2, _ = roc_curve(y_test, y_pred_l2_prob)

        # calculate the AUC
        roc_auc_l1 = auc(fpr_l1, tpr_l1)
        roc_auc_l2 = auc(fpr_l2, tpr_l2)

        plt.figure(figsize=(6, 4))

        # ROC L1
        plt.plot(fpr_l1, tpr_l1, color='blue', lw=2, label=f'L1 Regularization (A
        plt.fill_between(fpr_l1, tpr_l1, color='blue', alpha=0.2)

        # ROC L2
        plt.plot(fpr_l2, tpr_l2, color='green', lw=2, label=f'L2 Regularization (
        plt.fill_between(fpr_l2, tpr_l2, color='green', alpha=0.1)

        plt.plot([0, 1], [0, 1], color='gray', linestyle='--')


        plt.title('ROC Curve Comparison')
        plt.xlabel('False Positive Rate')
        plt.ylabel('True Positive Rate')

        plt.show()

[[6.92758810e-01 3.07241190e-01]
 [2.44841148e-07 9.9999755e-01]
 [6.03871558e-01 3.96128442e-01]]

[[6.77569240e-01 3.22430760e-01]
 [1.90215880e-05 9.99980978e-01]
 [5.65745607e-01 4.34254393e-01]]

```



#### L1 Regularization Model Report:

	precision	recall	f1-score	support
Survived	0.89	0.99	0.94	123
Death	0.99	0.84	0.90	91
accuracy			0.93	214
macro avg	0.94	0.91	0.92	214
weighted avg	0.93	0.93	0.92	214

#### L2 Regularization Model Report:


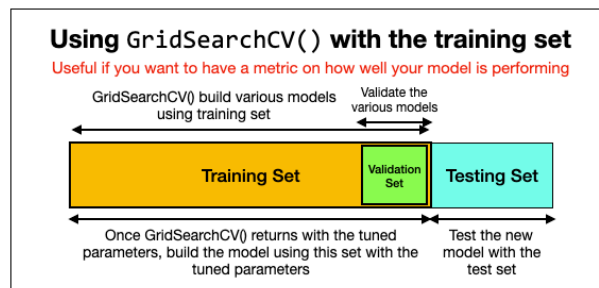
	precision	recall	f1-score	support
Survived	0.88	0.99	0.93	123
Death	0.99	0.82	0.90	91
accuracy			0.92	214
macro avg	0.94	0.91	0.92	214
weighted avg	0.93	0.92	0.92	214

So, for L1 is better **Choosing L1 Regularization**

## 4. Model Optimization

### 4.1 Hyperparameter Tuning

- **Hyperparameters** : are parameters that need to be manually set before training in machine learning models, such as the maximum depth of decision trees, the C value of support vector machines, etc.
- **Grid Search** : refers to traversing the specified hyperparameter space, training and validating each set of hyperparameter combinations, and finding the best hyperparameter combination.
- **Cross Validation (CV)** : is a technique for evaluating model performance, dividing the data into multiple subsets, each subset is used as a validation set in turn, and the rest are used as training sets, thereby reducing performance fluctuations caused by different data divisions



## Method 1: Grid Search (CV)

```
In [ ]: # Use 'GridSearchCV' or 'RandomizedSearchCV' to find the best hyperparameter
from sklearn.model_selection import GridSearchCV
# from sklearn.ensemble import RandomForestClassifier

# choosing model with L1 Regularization
# model_l1

param_grid = {
    'C': [0.01, 0.1, 1, 10, 100],
    'solver': ['liblinear', 'saga'],
    'penalty': ['l1', 'l2'],
    'max_iter': [1000],
    'tol': [1e-3, 1e-1],
    'class_weight': ['balanced'] # 自动调整类别权重
}

grid_search = GridSearchCV(estimator=model_l1, param_grid=param_grid, cv=5)
grid_search.fit(X_train, y_train)

print("Best hyper param:\n", grid_search.best_params_)
```

Best hyper param:

```
{'C': 0.1, 'class_weight': 'balanced', 'max_iter': 1000, 'penalty': 'l1',
'solver': 'saga', 'tol': 0.001}
```

```
In [ ]: # aggrisin the state_dict into the model
best_model = grid_search.best_estimator_
best_model
```

```
Out[ ]: ▼ LogisticRegression
LogisticRegression(C=0.1, class_weight='balanced', max_iter=1000,
penalty='l1',
random_state=42, solver='saga', tol=0.001)
```

```
In [ ]: y_pred_best = best_model.predict(X_test)

print('Best Model Report:\n', classification_report(
    y_test, y_pred_best,
    target_names=['Survived', 'Death'])
)

y_pred_best_prob = best_model.predict_proba(X_test)[: , 1]
# print(best_model.predict_proba(X_test)[:3])

fpr_best, tpr_best, _ = roc_curve(y_test, y_pred_best_prob)

# calculate the AUC
roc_auc_best = auc(fpr_best, tpr_best)
```

Best Model Report:

	precision	recall	f1-score	support
Survived	0.89	0.99	0.94	123
Death	0.99	0.84	0.90	91
accuracy			0.93	214
macro avg	0.94	0.91	0.92	214
weighted avg	0.93	0.93	0.92	214

```
In [ ]: # Evaluate the model with the best parameters on the test set

# Best Model
y_pred_best = best_model.predict(X_test)

print('Best Model Report (Grid Search):\n', classification_report(
    y_test, y_pred_best,
    target_names=['Survived', 'Death'])
)

print('----' * 20)
print()
y_pred_best_prob = best_model.predict_proba(X_test)[: , 1]
# print(best_model.predict_proba(X_test)[:3])

fpr_best, tpr_best, _ = roc_curve(y_test, y_pred_best_prob)

# calculate the AUC
roc_auc_best = auc(fpr_best, tpr_best)
```

```
#####

# Previous Model
print('Previous Model Report:\n', classification_report(
    y_test, y_pred_l2,
    target_names=['Survived', 'Death'])
)

# plotting
plt.figure(figsize=(5, 3))

# ROC L1
plt.plot(fpr_l1, tpr_l1, color='blue', lw=2, label=f'Previous Regularizat
plt.fill_between(fpr_l1, tpr_l1, color='blue', alpha=0.2)

plt.plot(fpr_best, tpr_best, color='red', lw=2, label=f'Best Regularizati
plt.fill_between(fpr_best, tpr_best, color='red', alpha=0.1)

plt.plot([0, 1], [0, 1], color='gray', linestyle='--')

plt.title('Comparision of Previous & Updated')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')


plt.show()
```

Best Model Report (Grid Search):

	precision	recall	f1-score	support
Survived	0.89	0.99	0.94	123
Death	0.99	0.84	0.90	91
accuracy			0.93	214
macro avg	0.94	0.91	0.92	214
weighted avg	0.93	0.93	0.92	214

Previous Model Report:

	precision	recall	f1-score	support
Survived	0.88	0.99	0.93	123
Death	0.99	0.82	0.90	91
accuracy			0.92	214
macro avg	0.94	0.91	0.92	214
weighted avg	0.93	0.92	0.92	214



## Method 2: RandomizedSearchCV

```
In [ ]: from sklearn.model_selection import RandomizedSearchCV
        from scipy.stats import uniform

        # choosing model with L1 Regularization
        # model_l1

        # Define the hyperparameter distribution for RandomizedSearchCV
        param_dist = {
            'C': uniform(0.01, 100), # Uniform distribution for C (from 0.01 to
            'solver': ['liblinear', 'saga'],
            'penalty': ['l1', 'l2'],
            'max_iter': [1000],
            'tol': uniform(1e-3, 1e-1), # Random distribution for tolerance
            'class_weight': ['balanced'] # Automatically adjust class weights
        }

        # Using RandomizedSearchCV with 5-fold cross-validation and n_jobs=-1 for
        random_search = RandomizedSearchCV(estimator=model_l1, param_distribution

        # Fit the model
        random_search.fit(X_train, y_train)

        # Print the best hyperparameters
        print("Best hyper parameters:\n", random_search.best_params_)
```

Best hyper parameters:

```
{'C': 72.83163486118596, 'class_weight': 'balanced', 'max_iter': 1000, 'p
enalty': 'l2', 'solver': 'liblinear', 'tol': 0.06423058305935796}
```

```
In [ ]: best_model1 = random_search.best_estimator_
        best_model1
```

Out [ ]:

```

▼ LogisticRegression
LogisticRegression(C=72.83163486118596, class_weight='balanced',
max_iter=1000,
                    random_state=42, solver='liblinear',
                    tol=0.06423058305935796)

```

In [ ]:

```

# Evaluate the model with the best parameters on the test set

# Best Model
y_pred_best1 = best_model1.predict(X_test)

print('Best Model Report (Random Search):\n', classification_report(
    y_test, y_pred_best1,
    target_names=['Survived', 'Death'])
)

print('---' * 20)
print()
y_pred_best_proba = best_model.predict_proba(X_test)[: , 1]
# print(best_model.predict_proba(X_test)[:3])

fpr_best1, tpr_best1, _ = roc_curve(y_test, y_pred_best_proba)

# calculate the AUC
roc_auc_best1 = auc(fpr_best1, tpr_best1)

#####

# Previous Model
print('Previous Model Report:\n', classification_report(
    y_test, y_pred_l1,
    target_names=['Survived', 'Death'])
)

# plotting
plt.figure(figsize=(5, 3))

# ROC Grid Search
plt.plot(fpr_best, tpr_best, color='blue', lw=2, label=f'Best Regularizat
plt.fill_between(fpr_best, tpr_best1, color='blue', alpha=0.1)

plt.plot(fpr_best1, tpr_best1, color='red', lw=2, label=f'Best Regulariza
plt.fill_between(fpr_best1, tpr_best1, color='red', alpha=0.1)

plt.plot([0, 1], [0, 1], color='gray', linestyle='--')

plt.title('Comparision of Previous & Updated')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')

plt.show()

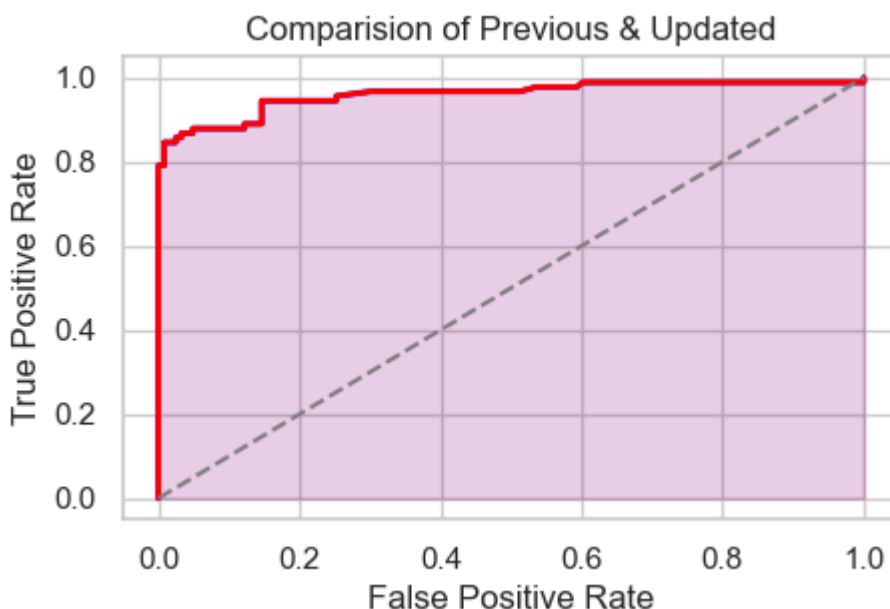
```

## Best Model Report (Random Search):

	precision	recall	f1-score	support
Survived	0.91	0.98	0.94	123
Death	0.96	0.87	0.91	91
accuracy			0.93	214
macro avg	0.94	0.92	0.93	214
weighted avg	0.93	0.93	0.93	214

## Previous Model Report:

	precision	recall	f1-score	support
Survived	0.89	0.99	0.94	123
Death	0.99	0.84	0.90	91
accuracy			0.93	214
macro avg	0.94	0.91	0.92	214
weighted avg	0.93	0.93	0.92	214



so, no obvious improvement for Random Search...

## 5. Model Deployment

### 5.1 Model Saving

```
In [ ]: import joblib

# save the model into a file
joblib.dump(best_model, 'titanic_analysis.pkl')
```

```
Out[ ]: ['titanic_analysis.pkl']
```



## 5.2 Model Deployment using Function to Implement

```
# python prediction.py titanic.csv titanic.pkl
predict('titanic.csv', 'titanic_analysis.pkl')
```

```
In [ ]: def predict(data_path, model_path):
        # pkl as surffix name
        model = joblib.load(model_path)

        # csv file input
        new_data = pd.read_csv(data_path)

        # make prediction
        predictions = model.predict(new_data)

        print(predictions)
```

```
In [ ]: # predict('titanic.csv', 'titanic_analysis.pkl')
```

## 5.3 Command Line Conduction

### prediction.py File

- Run Command By:

```
python prediction.py titanic.csv titanic.pkl
```

```
import argparse
import pandas as pd
import joblib

parser = argparse.ArgumentParser(description='Predict data
classification using a specified model')

parser.add_argument('input_file', type=str, help='Path to the CSV
file for prediction')
parser.add_argument('model_path', type=str, help='Path to the
trained model file')

args = parser.parse_args()

# load in command's input
model = joblib.load(args.model_path)
new_data = pd.read_csv(args.input_file)

predictions = model_l1.predict(new_data)

print("Predictions:", predictions)
```